

Importing Necessary Libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from statsmodels.tsa.arima.model import ARIMA
from statsmodels.tsa.statespace.sarimax import SARIMAX
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
from statsmodels.tsa.stattools import acf, pacf
from statsmodels.stats.diagnostic import het_arch
from statsmodels.tsa.stattools import adfuller
import pickle
import warnings
warnings.filterwarnings("ignore")
pd.options.display.float_format = '{:,.4f}'.format
from sklearn.preprocessing import LabelEncoder
```

Loading Dataset

```
train_data = pd.read_csv(r"C:\Users\asus\Downloads\train (1).csv")
oil_data = pd.read_csv(r"C:\Users\asus\Downloads\data\oil.csv")
```

Convert into Date-Time Format

```
train_data['date'] = pd.to_datetime(train_data['date'])
oil_data['date'] = pd.to_datetime(oil_data['date'])
```

Filling Missing Values

```
oil_data['dcoilwtico'] = oil_data['dcoilwtico'].fillna(method='bfill')
oil_data['dcoilwtico'] = oil_data['dcoilwtico'].fillna(method='ffill')
oil_data
```

	date	dcoilwtico
0	2020-01-01	93.1400
1	2020-01-02	93.1400
2	2020-01-03	92.9700
3	2020-01-04	93.1200
4	2020-01-07	93.2000
...
1049	2024-01-09	51.9500
1050	2024-01-10	50.8200
1051	2024-01-11	52.1900
1052	2024-01-12	53.0100
1053	2024-01-13	52.3600

```
[1054 rows x 2 columns]
```

Merge Train Data with Oil Data

```
merged_data = pd.merge(train_data, oil_data, on='date', how='left')
merged_data['family'] = merged_data['family'].replace('BREAD/BAKERY',
'BREAD_BAKERY')
merged_data['dcoilwtico'] =
merged_data['dcoilwtico'].fillna(method='bfill')
merged_data['dcoilwtico'] =
merged_data['dcoilwtico'].fillna(method='ffill')
merged_data
```

	id	date	store_nbr	family
sales \				
0	0	2020-01-01	1	AUTOMOTIVE
0.0000				
1	1	2020-01-01	1	BABY CARE
0.0000				
2	2	2020-01-01	1	BEAUTY
0.0000				
3	3	2020-01-01	1	BEVERAGES
0.0000				
4	4	2020-01-01	1	BOOKS
0.0000				
...
...				
2596369	2596369	2023-12-31	9	POULTRY
687.8530				
2596370	2596370	2023-12-31	9	PREPARED FOODS
100.4050				
2596371	2596371	2023-12-31	9	PRODUCE
3,091.3560				
2596372	2596372	2023-12-31	9	SCHOOL AND OFFICE SUPPLIES
2.0000				
2596373	2596373	2023-12-31	9	SEAFOOD
13.0000				

	onpromotion	dcoilwtico
0	0	93.1400
1	0	93.1400
2	0	93.1400
3	0	93.1400
4	0	93.1400
...
2596369	1	53.7500
2596370	1	53.7500
2596371	3	53.7500
2596372	0	53.7500
2596373	2	53.7500

[2596374 rows x 7 columns]

Exploratory Data Analysis

Checking Missing Values

```
merged_data.isnull().sum()
```

```
id          0
date        0
store_nbr   0
family      0
sales       0
onpromotion 0
dcoilwtico  0
dtype: int64
```

```
merged_data.drop(columns = ["id", "onpromotion"], inplace = True)
merged_data
```

	date	store_nbr	family	
sales \				
0	2020-01-01	1	AUTOMOTIVE	0.0000
1	2020-01-01	1	BABY CARE	0.0000
2	2020-01-01	1	BEAUTY	0.0000
3	2020-01-01	1	BEVERAGES	0.0000
4	2020-01-01	1	BOOKS	0.0000
...
2596369	2023-12-31	9	POULTRY	687.8530
2596370	2023-12-31	9	PREPARED FOODS	100.4050
2596371	2023-12-31	9	PRODUCE	3,091.3560
2596372	2023-12-31	9	SCHOOL AND OFFICE SUPPLIES	2.0000
2596373	2023-12-31	9	SEAFOOD	13.0000

	dcoilwtico
0	93.1400
1	93.1400
2	93.1400
3	93.1400
4	93.1400
...	...
2596369	53.7500

```
2596370    53.7500
2596371    53.7500
2596372    53.7500
2596373    53.7500
```

```
[2596374 rows x 5 columns]
```

```
print(merged_data.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2596374 entries, 0 to 2596373
Data columns (total 5 columns):
```

```
#    Column      Dtype
---  -
0    date      datetime64[ns]
1    store_nbr  int64
2    family     object
3    sales      float64
4    dcoilwtico float64
```

```
dtypes: datetime64[ns](1), float64(2), int64(1), object(1)
```

```
memory usage: 99.0+ MB
```

```
None
```

```
print(merged_data.describe())
```

	date	store_nbr	sales \
count	2596374	2,596,374.0000	2,596,374.0000
mean	2021-12-31 06:24:27.673301248	27.5000	338.7139
min	2020-01-01 00:00:00	1.0000	0.0000
25%	2021-01-01 00:00:00	14.0000	0.0000
50%	2022-01-01 00:00:00	27.5000	9.0000
75%	2023-01-01 00:00:00	41.0000	184.0608
max	2023-12-31 00:00:00	54.0000	124,717.0000
std	NaN	15.5858	1,055.8287

	dcoilwtico
count	2,596,374.0000
mean	70.7374
min	26.1900
25%	46.2100
50%	60.0100
75%	97.0000
max	110.6200
std	26.4728

```
print(merged_data['family'].nunique(), "unique product families")
```


```
33 unique product families
```

```
print(merged_data['store_nbr'].nunique(), "unique stores")
```

54 unique stores

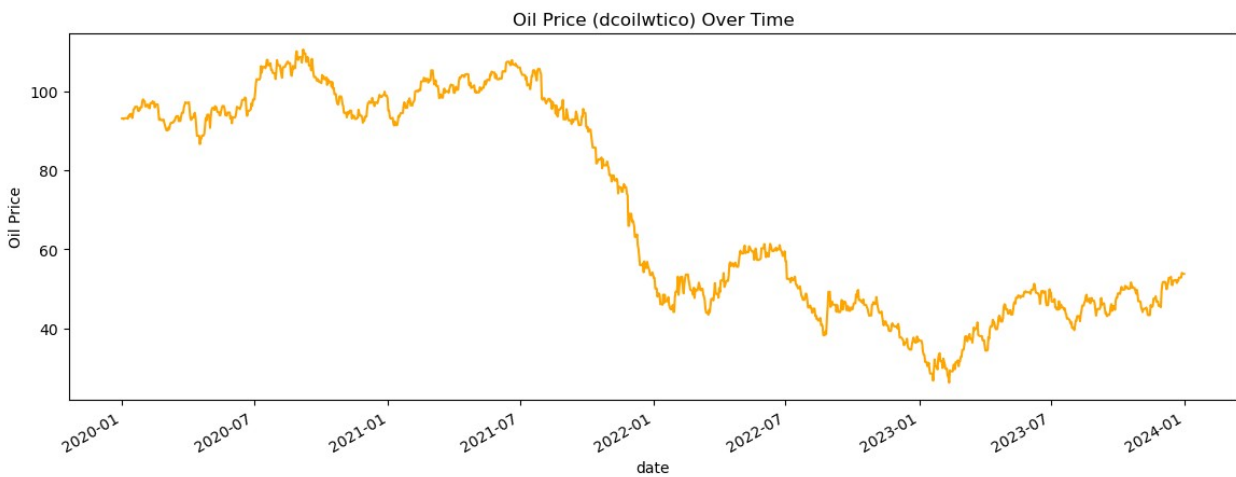
Sales Over Time (Total)

```
merged_data.groupby('date')['sales'].sum().plot(figsize=(14, 5))  
plt.title("Total Sales Over Time")  
plt.ylabel("Sales")  
plt.show()
```




Oil Price Over Time

```
merged_data.groupby('date')['dcoilwtico'].mean().plot(figsize=(14, 5),  
color='orange')  
plt.title("Oil Price (dcoilwtico) Over Time")  
plt.ylabel("Oil Price")  
plt.show()
```



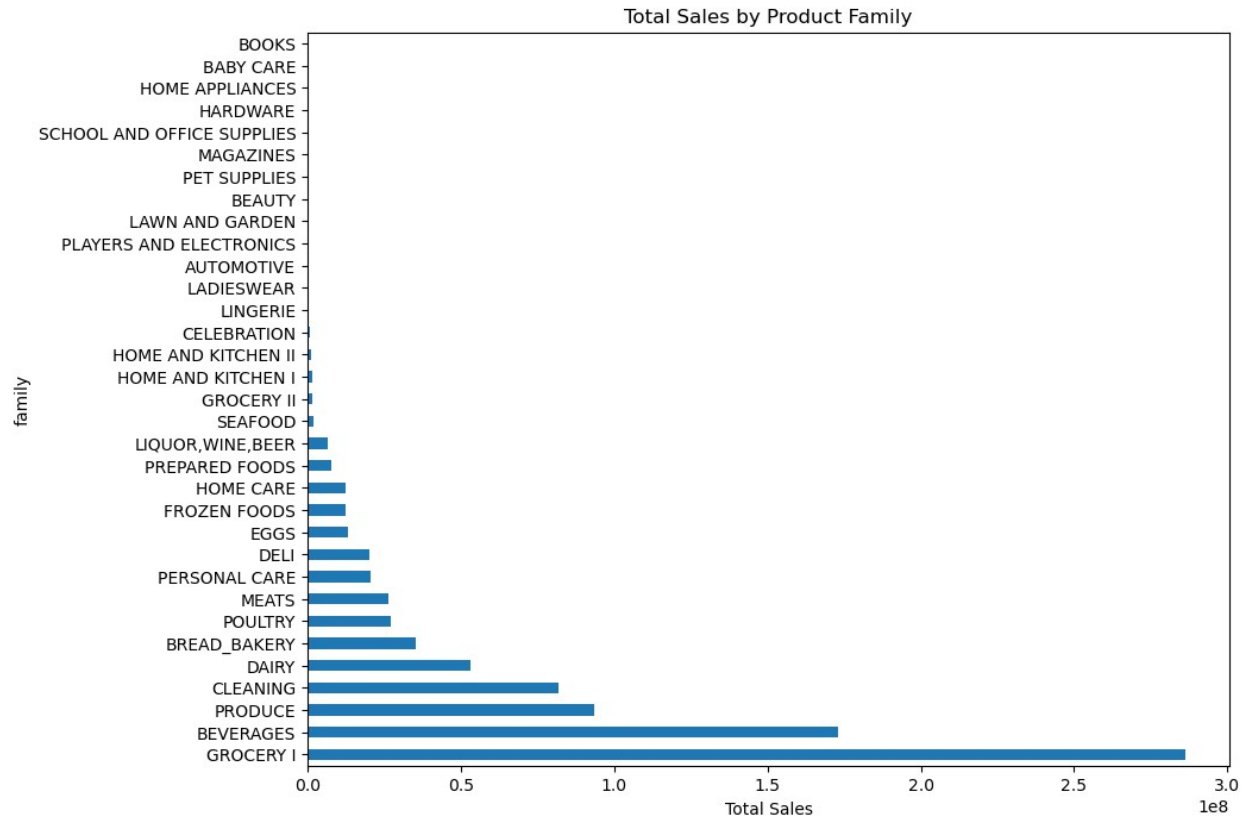
Sales by Store

```
store_sales = merged_data.groupby('store_nbr')  
['sales'].sum().sort_values(ascending=False)  
store_sales.plot(kind='bar', figsize=(10,5))  
plt.title("Total Sales by Store")  
plt.xlabel("Store Number")  
plt.ylabel("Total Sales")  
plt.show()
```



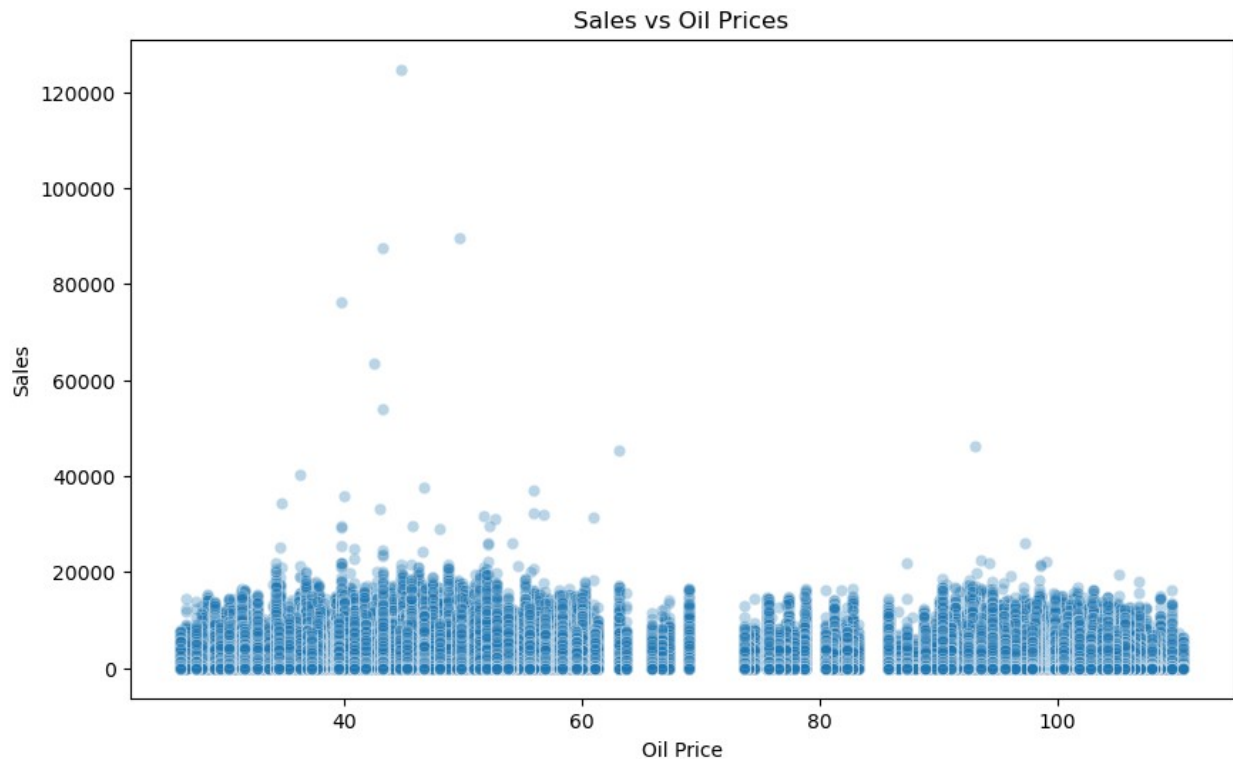
Sales by Product Family

```
family_sales = merged_data.groupby('family')  
['sales'].sum().sort_values(ascending=False)  
family_sales.plot(kind='barh', figsize=(10, 8))  
plt.title("Total Sales by Product Family")  
plt.xlabel("Total Sales")  
plt.show()
```



Relationship Between Sales & Oil Prices

```
plt.figure(figsize=(10,6))
sns.scatterplot(data=merged_data, x='dcoilwtico', y='sales',
alpha=0.3)
plt.title("Sales vs Oil Prices")
plt.xlabel("Oil Price")
plt.ylabel("Sales")
plt.show()
```



Correlation Matrix

```
merged_data[['sales', 'dcoilwtico']].corr()
```

	sales	dcoilwtico
sales	1.0000	-0.0733
dcoilwtico	-0.0733	1.0000

Splitting Date Yearly, Monthly, and Quarterly

```
merged_data["Year"] = merged_data["date"].dt.year
merged_data["Month"] = merged_data["date"].dt.month
merged_data["Quarter"] = merged_data["date"].dt.quarter
merged_data
```

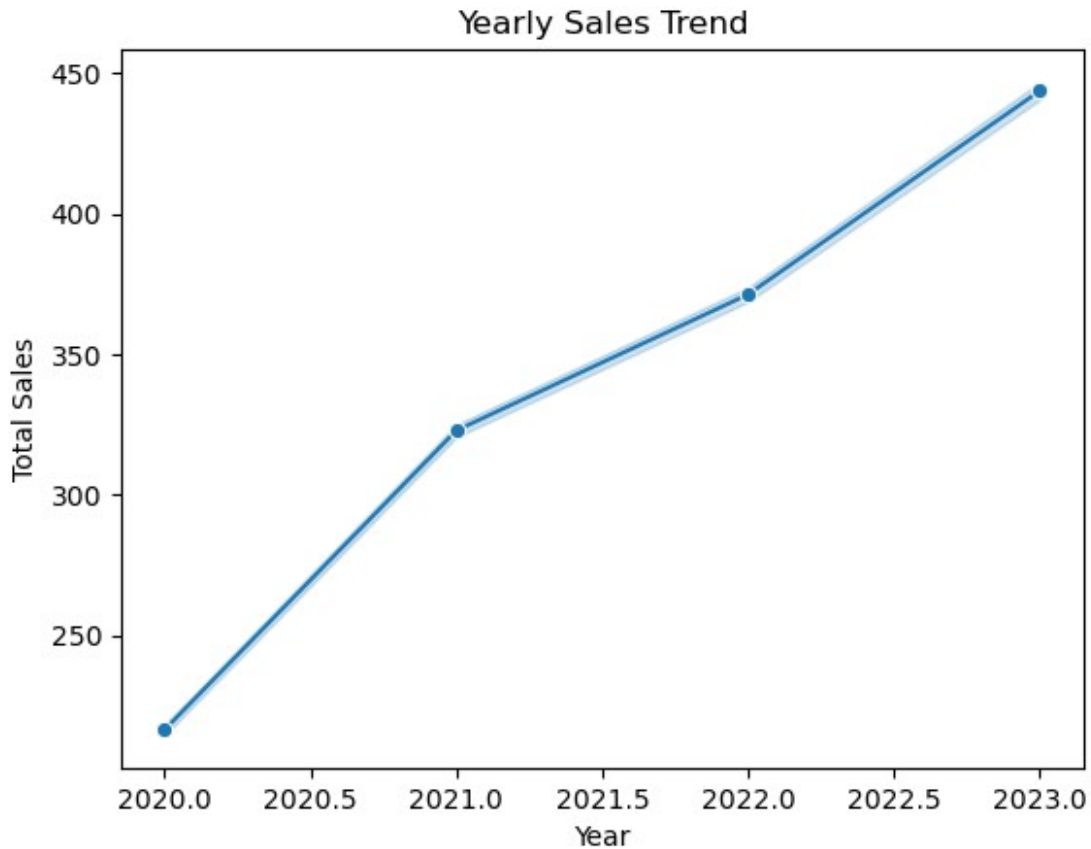
	date	store_nbr	family	
sales \				
0	2020-01-01	1	AUTOMOTIVE	0.0000
1	2020-01-01	1	BABY CARE	0.0000
2	2020-01-01	1	BEAUTY	0.0000
3	2020-01-01	1	BEVERAGES	0.0000
4	2020-01-01	1	BOOKS	0.0000

...
2596369	2023-12-31	9	POULTRY	687.8530
2596370	2023-12-31	9	PREPARED FOODS	100.4050
2596371	2023-12-31	9	PRODUCE	3,091.3560
2596372	2023-12-31	9	SCHOOL AND OFFICE SUPPLIES	2.0000
2596373	2023-12-31	9	SEAFOOD	13.0000

	dcoilwtico	Year	Month	Quarter
0	93.1400	2020	1	1
1	93.1400	2020	1	1
2	93.1400	2020	1	1
3	93.1400	2020	1	1
4	93.1400	2020	1	1
...
2596369	53.7500	2023	12	4
2596370	53.7500	2023	12	4
2596371	53.7500	2023	12	4
2596372	53.7500	2023	12	4
2596373	53.7500	2023	12	4

[2596374 rows x 8 columns]

```
sns.lineplot(data=merged_data, x='Year', y='sales', marker='o')
plt.title("Yearly Sales Trend")
plt.xlabel("Year")
plt.ylabel("Total Sales")
plt.show()
```



Monthly Sales

```
merged_data["Year-Month"] = merged_data["Year"].astype(str) + "-" +
merged_data["Month"].astype(str).str.zfill(2)
merged_data
```


	date	store_nbr	family	
sales \				
0	2020-01-01	1	AUTOMOTIVE	0.0000
1	2020-01-01	1	BABY CARE	0.0000
2	2020-01-01	1	BEAUTY	0.0000
3	2020-01-01	1	BEVERAGES	0.0000
4	2020-01-01	1	BOOKS	0.0000
...
2596369	2023-12-31	9	POULTRY	687.8530
2596370	2023-12-31	9	PREPARED FOODS	100.4050

2596371	2023-12-31	9		PRODUCE	3,091.3560
2596372	2023-12-31	9	SCHOOL AND OFFICE SUPPLIES		2.0000
2596373	2023-12-31	9		SEAFOOD	13.0000

	dcoilwtico	Year	Month	Quarter	Year-Month
0	93.1400	2020	1	1	2020-01
1	93.1400	2020	1	1	2020-01
2	93.1400	2020	1	1	2020-01
3	93.1400	2020	1	1	2020-01
4	93.1400	2020	1	1	2020-01
...
2596369	53.7500	2023	12	4	2023-12
2596370	53.7500	2023	12	4	2023-12
2596371	53.7500	2023	12	4	2023-12
2596372	53.7500	2023	12	4	2023-12
2596373	53.7500	2023	12	4	2023-12

[2596374 rows x 9 columns]

```
plt.figure(figsize = (20,5))
sns.lineplot(data=merged_data, x='Year-Month', y='sales', marker='o')
plt.title("Monthly Sales Trend")
plt.xlabel("Year-Month")
plt.ylabel("Total Sales")
plt.xticks(rotation=45)
plt.show()
```



Quarterly Sales

```
merged_data["Year-Quarter"] = merged_data["Year"].astype(str) + "-" +
merged_data["Quarter"].astype(str).str.zfill(2)
merged_data
```

	date	store_nbr	family
sales \			

0	2020-01-01	1	AUTOMOTIVE	0.0000
1	2020-01-01	1	BABY CARE	0.0000
2	2020-01-01	1	BEAUTY	0.0000
3	2020-01-01	1	BEVERAGES	0.0000
4	2020-01-01	1	BOOKS	0.0000
...
2596369	2023-12-31	9	POULTRY	687.8530
2596370	2023-12-31	9	PREPARED FOODS	100.4050
2596371	2023-12-31	9	PRODUCE	3,091.3560
2596372	2023-12-31	9	SCHOOL AND OFFICE SUPPLIES	2.0000
2596373	2023-12-31	9	SEAFOOD	13.0000

	dcoilwtico	Year	Month	Quarter	Year-Month	Year-Quarter
0	93.1400	2020	1	1	2020-01	2020-01
1	93.1400	2020	1	1	2020-01	2020-01
2	93.1400	2020	1	1	2020-01	2020-01
3	93.1400	2020	1	1	2020-01	2020-01
4	93.1400	2020	1	1	2020-01	2020-01
...
2596369	53.7500	2023	12	4	2023-12	2023-04
2596370	53.7500	2023	12	4	2023-12	2023-04
2596371	53.7500	2023	12	4	2023-12	2023-04
2596372	53.7500	2023	12	4	2023-12	2023-04
2596373	53.7500	2023	12	4	2023-12	2023-04

[2596374 rows x 10 columns]

```
sns.lineplot(data=merged_data, x='Year-Quarter', y='sales',
marker='o')
plt.title("Quarterly Sales Trend")
plt.xlabel("Year-Quarter")
plt.ylabel("Total Sales")
plt.xticks(rotation=45)
plt.show()
```



Group By:

```
df = merged_data.groupby(["Year", "Quarter", "family"], as_index =
False).agg({"sales" : "sum", "dcoilwtico" : "mean" })
df
```

	Year	Quarter	family	sales
dcoilwtico				
0	2020	1	AUTOMOTIVE	21,412.0000
94.4666				
1	2020	1	BABY CARE	0.0000
94.4666				
2	2020	1	BEAUTY	11,667.0000
94.4666				
3	2020	1	BEVERAGES	5,001,127.0000
94.4666				
4	2020	1	BOOKS	0.0000
94.4666				
..
...				
523	2023	4	POULTRY	1,966,067.5634

```

49.1047
524 2023      4      PREPARED FOODS      516,604.2458
49.1047
525 2023      4      PRODUCE 11,354,309.3796
49.1047
526 2023      4  SCHOOL AND OFFICE SUPPLIES      7,721.0000
49.1047
527 2023      4      SEAFOOD      107,702.5450
49.1047

[528 rows x 5 columns]

```

Sort Index Quarterly

```

df['quarter_start'] = pd.PeriodIndex(year=df['Year'],
quarter=df['Quarter'], freq='Q').start_time

df.set_index('quarter_start', inplace=True)

df.sort_index(inplace=True)
df

```

	Year	Quarter	family
sales \ quarter_start			
2020-01-01	2020	1	AUTOMOTIVE
21,412.0000			
2020-01-01	2020	1	BABY CARE
0.0000			
2020-01-01	2020	1	BEAUTY
11,667.0000			
2020-01-01	2020	1	BEVERAGES
5,001,127.0000			
2020-01-01	2020	1	BOOKS
0.0000			
...
..			
2023-10-01	2023	4	POULTRY
1,966,067.5634			
2023-10-01	2023	4	PREPARED FOODS
516,604.2458			
2023-10-01	2023	4	PRODUCE
11,354,309.3796			
2023-10-01	2023	4	SCHOOL AND OFFICE SUPPLIES
7,721.0000			
2023-10-01	2023	4	SEAFOOD
107,702.5450			
	dcoilwtico		

quarter_start	
2020-01-01	94.4666
2020-01-01	94.4666
2020-01-01	94.4666
2020-01-01	94.4666
2020-01-01	94.4666
...	...
2023-10-01	49.1047
2023-10-01	49.1047
2023-10-01	49.1047
2023-10-01	49.1047
2023-10-01	49.1047

[528 rows x 5 columns]

```
# Initialize LabelEncoder
le = LabelEncoder()
```

```
# Fit and transform the 'family' column
df['family_encoded'] = le.fit_transform(df['family'])
```

```
# Create a mapping dictionary
family_mapping = dict(zip(le.classes_, le.transform(le.classes_)))
```

```
# Display the mapping
print("Category to Label Mapping:")
for category, label in family_mapping.items():
    print(f"{category}: {label}")
```

df

```
Category to Label Mapping:
AUTOMOTIVE: 0
BABY CARE: 1
BEAUTY: 2
BEVERAGES: 3
BOOKS: 4
BREAD_BAKERY: 5
CELEBRATION: 6
CLEANING: 7
DAIRY: 8
DELI: 9
EGGS: 10
FROZEN FOODS: 11
GROCERY I: 12
GROCERY II: 13
HARDWARE: 14
HOME AND KITCHEN I: 15
HOME AND KITCHEN II: 16
HOME APPLIANCES: 17
```

HOME CARE: 18
 LADIESWEAR: 19
 LAWN AND GARDEN: 20
 LINGERIE: 21
 LIQUOR,WINE,BEER: 22
 MAGAZINES: 23
 MEATS: 24
 PERSONAL CARE: 25
 PET SUPPLIES: 26
 PLAYERS AND ELECTRONICS: 27
 POULTRY: 28
 PREPARED FOODS: 29
 PRODUCE: 30
 SCHOOL AND OFFICE SUPPLIES: 31
 SEAFOOD: 32

sales \ quarter_start	Year	Quarter	family
2020-01-01 21,412.0000	2020	1	AUTOMOTIVE
2020-01-01 0.0000	2020	1	BABY CARE
2020-01-01 11,667.0000	2020	1	BEAUTY
2020-01-01 5,001,127.0000	2020	1	BEVERAGES
2020-01-01 0.0000	2020	1	BOOKS
...
..			
2023-10-01 1,966,067.5634	2023	4	POULTRY
2023-10-01 516,604.2458	2023	4	PREPARED FOODS
2023-10-01 11,354,309.3796	2023	4	PRODUCE
2023-10-01 7,721.0000	2023	4	SCHOOL AND OFFICE SUPPLIES
2023-10-01 107,702.5450	2023	4	SEAFOOD

quarter_start	dcoilwtico	family_encoded
2020-01-01	94.4666	0
2020-01-01	94.4666	1
2020-01-01	94.4666	2
2020-01-01	94.4666	3
2020-01-01	94.4666	4


```

...
2023-10-01      49.1047      28
2023-10-01      49.1047      29
2023-10-01      49.1047      30
2023-10-01      49.1047      31
2023-10-01      49.1047      32

```

```
[528 rows x 6 columns]
```

```
df.drop("family", axis = 1, inplace = True)
df
```

```
df.to_csv("Cleaned_Data.csv")
```

Differencing Data Quarterly

```
df["sales_lag"] = df["sales"].diff(33)
df
```

family_encoded \ quarter_start	Year	Quarter	sales	dcoilwtico
2020-01-01	2020	1	21,412.0000	94.4666
2020-01-01	2020	1	0.0000	94.4666
2020-01-01	2020	1	11,667.0000	94.4666
2020-01-01	2020	1	5,001,127.0000	94.4666
2020-01-01	2020	1	0.0000	94.4666
...
2023-10-01	2023	4	1,966,067.5634	49.1047
2023-10-01	2023	4	516,604.2458	49.1047
2023-10-01	2023	4	11,354,309.3796	49.1047
2023-10-01	2023	4	7,721.0000	49.1047
2023-10-01	2023	4	107,702.5450	49.1047

quarter_start	sales_lag
2020-01-01	NaN
2020-01-01	NaN

```

2020-01-01      NaN
2020-01-01      NaN
2020-01-01      NaN
...
2023-10-01    143,668.9392
2023-10-01      7,664.2459
2023-10-01    827,308.8577
2023-10-01    -34,795.0000
2023-10-01    -10,243.9800

[528 rows x 6 columns]

```

Dropping Null Values

```

df.dropna(inplace = True)
df

```

```

      family_encoded  Year  Quarter      sales  dcoilwtico
quarter_start
2020-04-01    2020      2    24,830.0000    94.1429
0
2020-04-01    2020      2      0.0000    94.1429
1
2020-04-01    2020      2    11,228.0000    94.1429
2
2020-04-01    2020      2  5,396,350.0000    94.1429
3
2020-04-01    2020      2      0.0000    94.1429
4
...
2023-10-01    2023      4  1,966,067.5634    49.1047
28
2023-10-01    2023      4    516,604.2458    49.1047
29
2023-10-01    2023      4  11,354,309.3796    49.1047
30
2023-10-01    2023      4      7,721.0000    49.1047
31
2023-10-01    2023      4    107,702.5450    49.1047
32

      sales_lag
quarter_start
2020-04-01    3,418.0000
2020-04-01      0.0000
2020-04-01   -439.0000
2020-04-01  395,223.0000

```

```

2020-04-01      0.0000
...
2023-10-01    143,668.9392
2023-10-01      7,664.2459
2023-10-01    827,308.8577
2023-10-01    -34,795.0000
2023-10-01   -10,243.9800

```

```
[495 rows x 6 columns]
```

```

df["Year-Quarter"] = df["Year"].astype(str) + "-" +
df["Quarter"].astype(str).str.zfill(2)
df

```

family_encoded \ quarter_start	Year	Quarter	sales	dcoilwtico
2020-04-01	2020	2	24,830.0000	94.1429
0				
2020-04-01	2020	2	0.0000	94.1429
1				
2020-04-01	2020	2	11,228.0000	94.1429
2				
2020-04-01	2020	2	5,396,350.0000	94.1429
3				
2020-04-01	2020	2	0.0000	94.1429
4				
...
..				
2023-10-01	2023	4	1,966,067.5634	49.1047
28				
2023-10-01	2023	4	516,604.2458	49.1047
29				
2023-10-01	2023	4	11,354,309.3796	49.1047
30				
2023-10-01	2023	4	7,721.0000	49.1047
31				
2023-10-01	2023	4	107,702.5450	49.1047
32				

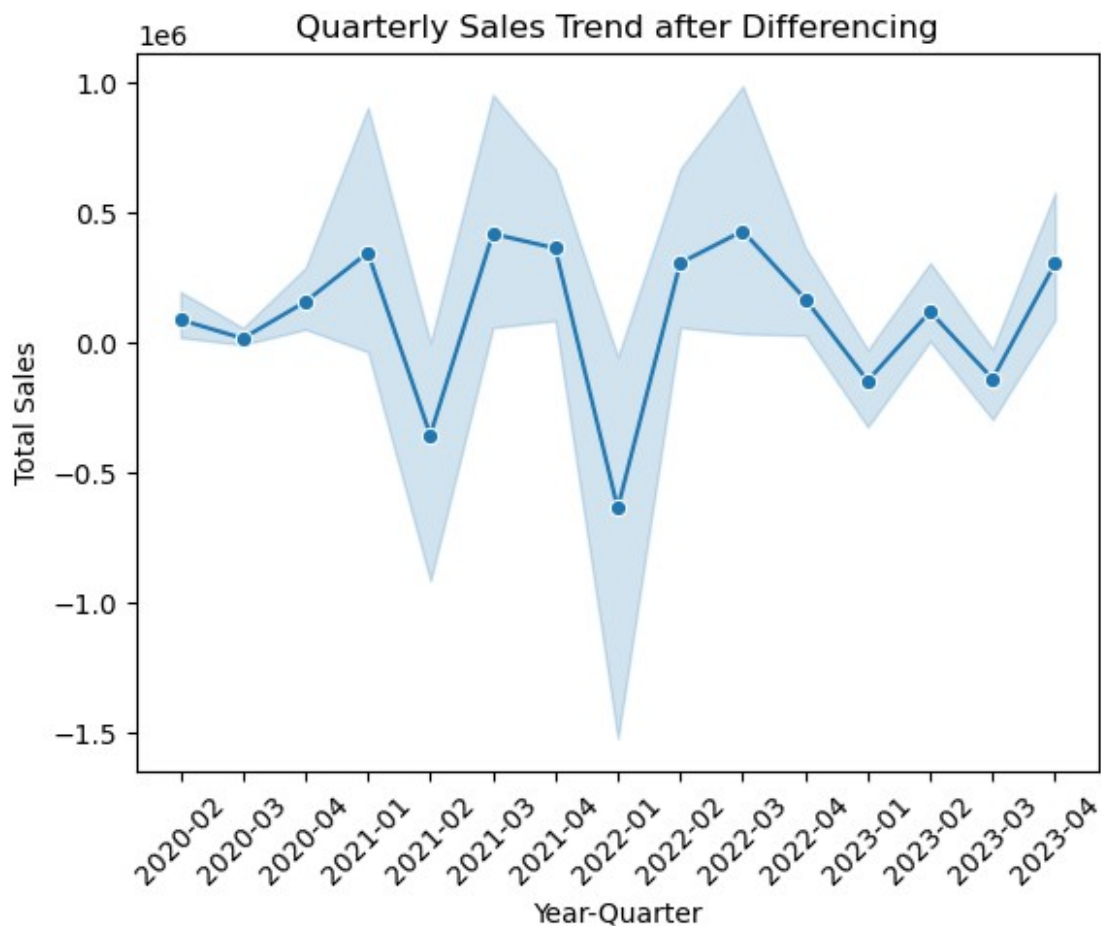
quarter_start	sales_lag	Year-Quarter
2020-04-01	3,418.0000	2020-02
2020-04-01	0.0000	2020-02
2020-04-01	-439.0000	2020-02
2020-04-01	395,223.0000	2020-02
2020-04-01	0.0000	2020-02
...
2023-10-01	143,668.9392	2023-04

2023-10-01	7,664.2459	2023-04
2023-10-01	827,308.8577	2023-04
2023-10-01	-34,795.0000	2023-04
2023-10-01	-10,243.9800	2023-04

[495 rows x 7 columns]

Checking Seasonality After Differencing

```
sns.lineplot(data=df, x='Year-Quarter', y='sales_lag', marker='o')
plt.title("Quarterly Sales Trend after Differencing")
plt.xlabel("Year-Quarter")
plt.ylabel("Total Sales")
plt.xticks(rotation=45)
plt.show()
```



```
df.to_csv("Differencing_1.csv")
```

Checking Stationarity using Augmented Dickey-Fuller Test

```
def check_stationarity(timeseries):
    result = adfuller(timeseries)
    print("ADF Statistic:", result[0])
    print("p-value:", result[1])
    print("Critical Values:", result[4])

    if result[1] < 0.05:
        print("The time series is stationary.")
    else:
        print("The time series is non-stationary.")

# Example usage
check_stationarity(df['sales_lag'])

ADF Statistic: -9.966937012191245
p-value: 2.2937883409570914e-17
Critical Values: {'1%': -3.4437936797256317, '5%': -2.867468682890213, '10%': -2.5699277594606915}
The time series is stationary.
```

Checking Homoscedasticity using ARCH

```
# Perform the ARCH test
arch_test = het_arch(df['sales_lag'])

print("ARCH Test Statistic:", arch_test[0])
print("p-value:", arch_test[1])

# Interpretation
if arch_test[1] < 0.05:
    print("Heteroscedasticity detected")
else:
    print("Homoscedasticity detected")

ARCH Test Statistic: 13.987843793358284
p-value: 0.17354681263428837
Homoscedasticity detected
```

Dropping rows containing INF and NaN Values

```
df.replace([np.inf, -np.inf], np.nan, inplace=True)
df.dropna(inplace=True)
df.isnull().sum()
```

Year	0
Quarter	0
sales	0
dcoilwtico	0

```
family_encoded    0
sales_lag         0
Year-Quarter      0
dtype: int64
```

```
df.shape
```

```
(495, 7)
```

```
df.drop(columns = ["sales", "Year-Quarter"], inplace = True)
df
```

	Year	Quarter	dcoilwtico	family_encoded	sales_lag
quarter_start					
2020-04-01	2020	2	94.1429	0	3,418.0000
2020-04-01	2020	2	94.1429	1	0.0000
2020-04-01	2020	2	94.1429	2	-439.0000
2020-04-01	2020	2	94.1429	3	395,223.0000
2020-04-01	2020	2	94.1429	4	0.0000
...
2023-10-01	2023	4	49.1047	28	143,668.9392
2023-10-01	2023	4	49.1047	29	7,664.2459
2023-10-01	2023	4	49.1047	30	827,308.8577
2023-10-01	2023	4	49.1047	31	-34,795.0000
2023-10-01	2023	4	49.1047	32	-10,243.9800

```
[495 rows x 5 columns]
```

Splitting Data for Training and Testing

```
train_size = int(0.8*len(df))
X_train = df.iloc[:train_size].copy()
Y_train = X_train.pop("sales_lag")
```

```
X_test = df.iloc[train_size:].copy()
Y_test = X_test.pop("sales_lag")
```

```
X_train.shape
```

```
(396, 4)
```

```
X_test.shape
```

```
(99, 4)
```

Getting ACF and PACF values

```
target = df['sales_lag']
```

```
max_lag = 33
```

```

acf_vals = acf(target, nlags=max_lag)
pacf_vals = pacf(target, nlags=max_lag)

lags = np.arange(max_lag + 1)
acf_pacf_temp_monthly = pd.DataFrame({'Lag': lags, 'ACF': acf_vals,
                                       'PACF': pacf_vals})

print(acf_pacf_temp_monthly)

```

	Lag	ACF	PACF
0	0	1.0000	1.0000
1	1	0.0252	0.0252
2	2	0.0096	0.0090
3	3	0.0041	0.0037
4	4	0.0231	0.0230
5	5	0.0212	0.0203
6	6	-0.1451	-0.1485
7	7	-0.0030	0.0042
8	8	-0.0002	0.0023
9	9	0.0655	0.0682
10	10	-0.0038	-0.0013
11	11	-0.0033	0.0018
12	12	0.0730	0.0534
13	13	-0.0001	-0.0068
14	14	0.0107	0.0071
15	15	0.0894	0.1138
16	16	-0.0010	-0.0109
17	17	-0.0084	-0.0147
18	18	0.0428	0.0601
19	19	0.0053	-0.0027
20	20	-0.0080	-0.0121
21	21	-0.0433	-0.0244
22	22	0.0097	0.0109
23	23	0.0101	0.0043
24	24	-0.0258	-0.0331
25	25	0.0046	0.0122
26	26	-0.0052	-0.0056
27	27	0.3929	0.4024
28	28	-0.0069	-0.0329
29	29	-0.0165	-0.0269
30	30	-0.0093	-0.0403
31	31	-0.0055	-0.0219
32	32	-0.0186	-0.0433
33	33	-0.3609	-0.3397

Plotting ACF and PACF

```

fig, axes = plt.subplots(nrows=1, ncols=2, figsize=(12, 4))
plot_acf(df['sales_lag'], ax=axes[0], lags=20)

```


```

axes[0].set_title('Autocorrelation Function (ACF)')

plot_pacf(df['sales_lag'], ax=axes[1], lags=20)
axes[1].set_title('Partial Autocorrelation Function (PACF)')

plt.tight_layout()
plt.show()

```



```

df['sales_lag'].describe()

count          495.0000
mean           94,743.9112
std            1,045,943.2266
min            -10,250,091.0875
25%            -1,827.5000
50%             1,244.0000
75%            39,661.5496
max             7,207,902.7638
Name: sales_lag, dtype: float64

```

Applying SARIMAX

```

model = SARIMAX(Y_train, exog=X_train, order=(6, 0, 6)) # adjust
(p,d,q) based on AIC/BIC/grid search
model_fit = model.fit(dis=False)

# Forecast
pred = model_fit.predict(start=len(Y_train), end=len(Y_train)
+len(Y_test)-1, exog=X_test)

# Optional: evaluate model
from sklearn.metrics import mean_squared_error
import numpy as np

```



```
rmse = np.sqrt(mean_squared_error(Y_test, pred))
print("RMSE:", rmse)
```

RMSE: 639050.4538762666

Applying SARIMAX by Product Family

```
rms = {}
error = []
adequacy = {}
adequacy["Max"] = []
adequacy["Min"] = []
adequacy["RMSE"] = []
adequacy["RMSE_Ratio"] = []

for family in df['family_encoded'].unique():
    try:
        temp_df = df[df["family_encoded"] == family]
        train_value = int(0.8*len(temp_df))
        X_train = temp_df.iloc[:train_value].copy()
        Y_train = X_train.pop("sales_lag")

        model = SARIMAX(Y_train, exog=X_train, order=(6, 1, 6),
enforce_stationarity=False) # adjust (p,d,q) based on AIC/BIC/grid
search
        model_fit = model.fit()

        X_test = temp_df.iloc[train_value:].copy()
        Y_test = X_test.pop("sales_lag")

        Y_pred = model_fit.predict(start=len(Y_train),
end=len(Y_train)+len(Y_test)-1, exog=X_test)

        rmse = np.sqrt(mean_squared_error(Y_pred, Y_test))

        mean_sales = temp_df['sales_lag'].mean()
        rmse_ratio = rmse / mean_sales
        adequacy["RMSE_Ratio"].append(rmse_ratio)
        adequacy["Max"].append(temp_df['sales_lag'].max())
        adequacy["Min"].append(temp_df['sales_lag'].min())

        print(f"{family}done...")
        adequacy["RMSE"].append(rmse)
```

```

except:
    error.append(family)
pass

```

```

0done...
1done...
2done...
3done...
4done...
5done...
6done...
7done...
8done...
9done...
10done...
11done...
12done...
13done...
14done...
15done...
16done...
17done...
18done...
19done...
20done...
21done...
22done...
23done...
24done...
25done...
26done...
27done...
28done...
29done...
30done...
31done...
32done...

```

Checking Model Adequacy

```

data_frame = pd.DataFrame(adequacy)
data_frame

```

	Max	Min	RMSE	RMSE_Ratio
0	6,517.0000	-2,956.0000	2,831.3890	3.3879
1	926.0000	-832.0000	521.3647	5.5347
2	5,901.0000	-3,937.0000	4,222.7258	3.5445
3	4,543,100.0000	-6,784,893.0000	4,387,117.5931	5.5506
4	4,119.0000	0.0000	2,378.1058	8.6603

5	430,127.1401	-150,807.3087	123,876.6208	1.9591
6	48,801.0000	-61,302.0000	34,977.4479	7.0603
7	565,941.0000	-362,573.0000	615,558.6632	5.3067
8	895,388.0000	-217,301.0000	284,961.2107	1.4804
9	155,230.6540	-131,518.0310	203,463.7953	4.8286
10	127,991.0000	-40,371.0000	94,355.2366	3.5901
11	1,057,098.4722	-1,004,931.3481	584,755.5863	7.0356
12	3,870,652.8930	-2,444,630.5180	2,953,465.1745	3.9334
13	40,358.0000	-28,619.0000	23,102.2714	8.4178
14	1,410.0000	-1,902.0000	672.5326	6.5210
15	81,798.0000	-70,797.0000	18,982.3664	1.6737
16	55,888.0000	-45,509.0000	189,965.9074	21.1559
17	1,680.0000	-1,195.0000	1,794.8607	-47.7357
18	869,724.0000	-1,286,497.0000	1,063,601.7754	10.8655
19	41,950.0000	-64,004.0000	55,798.7126	15.0938
20	17,138.0000	-1,380.0000	9,543.4334	4.5416
21	7,735.0000	-8,956.0000	2,503.9843	-2.3149
22	152,267.0000	-135,817.0000	75,990.2606	3.7127
23	24,726.0000	-6,735.0000	16,921.0241	7.5435
24	115,918.1619	-156,887.7559	72,864.1260	3.4364
25	161,177.0000	-208,309.0000	168,399.3446	3.8279
26	15,513.0000	-27,036.0000	20,379.6217	8.1174
27	38,800.0000	-47,347.0000	22,074.4415	5.4822
28	612,834.3281	-133,655.1271	240,732.5353	3.4110
29	50,420.7129	-15,584.4129	4,490.5033	0.6620
30	7,207,902.7638	-10,250,091.0875	9,091,844.1498	12.0154
31	31,418.0000	-34,795.0000	27,636.4915	53.6909
32	26,342.4306	-12,322.4981	4,674.9120	-113.7439

Predicting Sales by Product Family

```

predictions = [] # to store results

for family in df['family_encoded'].unique():
    try:
        temp_df = df[df["family_encoded"] == family]
        train_value = int(0.8 * len(temp_df))
        X_train = temp_df.iloc[:train_value].copy()
        Y_train = X_train.pop("sales_lag")

        model = SARIMAX(Y_train, exog=X_train, order=(6, 1, 6),
enforce_stationarity=False)
        model_fit = model.fit(dispatch=False)

        X_test = temp_df.iloc[train_value:].copy()
        Y_test = X_test.pop("sales_lag")

        Y_pred = model_fit.predict(
            start=len(Y_train),

```

```

        end=len(Y_train) + len(Y_test) - 1,
        exog=X_test
    )

    # Extract corresponding date info from X_test
    test_dates = X_test.index

    for date, pred in zip(test_dates, Y_pred):
        year = pd.to_datetime(date).year
        quarter = pd.to_datetime(date).quarter
        predictions.append({
            'family': family,
            'year': year,
            'quarter': quarter,
            'prediction': pred
        })

    print(f"{family} done...")

except Exception as e:
    error.append(family)
    print(f"Error in {family}: {e}")
    pass

```

```

0 done...
1 done...
2 done...
3 done...
4 done...
5 done...
6 done...
7 done...
8 done...
9 done...
10 done...
11 done...
12 done...
13 done...
14 done...
15 done...
16 done...
17 done...
18 done...
19 done...
20 done...
21 done...
22 done...
23 done...
24 done...
25 done...

```

```
26 done...
27 done...
28 done...
29 done...
30 done...
31 done...
32 done...
```

Creating Data Frame for Predicted Sales

```
pred_df = pd.DataFrame(predictions)
pred_df
```

	family	year	quarter	prediction
0	0	2023	2	-803.5572
1	0	2023	3	-755.9646
2	0	2023	4	-1,676.1854
3	1	2023	2	708.8533
4	1	2023	3	591.0062
...
94	31	2023	3	17,939.5781
95	31	2023	4	12,829.6736
96	32	2023	2	-2,223.7848
97	32	2023	3	-5,462.5529
98	32	2023	4	-7,847.4777

```
[99 rows x 4 columns]
```

Sorting Predicted Sales Quarterly

```
# Step 1: Keep only the required columns
df_filtered = pred_df[['year', 'quarter', 'prediction',
                       'family']].copy()

# Step 2: Create quarter_start as a datetime index
df_filtered['quarter_start'] = pd.to_datetime(
    df_filtered['year'].astype(str) + 'Q' +
    df_filtered['quarter'].astype(str)
).dt.to_period('Q').dt.start_time

# Step 3: Sort by Quarter (2, then 3, then 4), then by Family
quarter_order = [2, 3, 4]
df_filtered['quarter'] = pd.Categorical(df_filtered['quarter'],
categories=quarter_order, ordered=True)
df_sorted = df_filtered.sort_values(by=['quarter', 'family'])

# Step 4: Set quarter_start as index and keep only the desired columns
df_final = df_sorted.set_index('quarter_start')[['year', 'quarter',
'family', 'prediction']]
```

```
df_final.columns = ['Year', 'Quarter', 'family', 'Prediction']
```

```
df_final
```

	Year	Quarter	family	Prediction
quarter_start				
2023-04-01	2023	2	0	-803.5572
2023-04-01	2023	2	1	708.8533
2023-04-01	2023	2	2	5,080.0664
2023-04-01	2023	2	3	3,263,638.9535
2023-04-01	2023	2	4	0.0000
...
2023-10-01	2023	4	28	-155,025.3304
2023-10-01	2023	4	29	3,438.8906
2023-10-01	2023	4	30	11,804,841.5150
2023-10-01	2023	4	31	12,829.6736
2023-10-01	2023	4	32	-7,847.4777

```
[99 rows x 4 columns]
```

```
df["Year-Quarter"] = df["Year"].astype(str) + "-" +  
df["Quarter"].astype(str).str.zfill(2)  
df
```

	Year	Quarter	dcoilwtico	family_encoded	sales_lag
\ quarter_start					
2020-04-01	2020	2	94.1429	0	3,418.0000
2020-04-01	2020	2	94.1429	1	0.0000
2020-04-01	2020	2	94.1429	2	-439.0000
2020-04-01	2020	2	94.1429	3	395,223.0000
2020-04-01	2020	2	94.1429	4	0.0000
...
2023-10-01	2023	4	49.1047	28	143,668.9392
2023-10-01	2023	4	49.1047	29	7,664.2459
2023-10-01	2023	4	49.1047	30	827,308.8577
2023-10-01	2023	4	49.1047	31	-34,795.0000
2023-10-01	2023	4	49.1047	32	-10,243.9800

quarter_start	Year-Quarter
2020-04-01	2020-02
2020-04-01	2020-02
2020-04-01	2020-02
2020-04-01	2020-02
2020-04-01	2020-02
...	...
2023-10-01	2023-04
2023-10-01	2023-04
2023-10-01	2023-04
2023-10-01	2023-04
2023-10-01	2023-04

[495 rows x 6 columns]

```
df_final["Year-Quarter"] = df_final["Year"].astype(str) + "-" +
df_final["Quarter"].astype(str).str.zfill(2)
df_final
```

quarter_start	Year	Quarter	family	Prediction	Year-Quarter
2023-04-01	2023	2	0	-803.5572	2023-02
2023-04-01	2023	2	1	708.8533	2023-02
2023-04-01	2023	2	2	5,080.0664	2023-02
2023-04-01	2023	2	3	3,263,638.9535	2023-02
2023-04-01	2023	2	4	0.0000	2023-02
...
2023-10-01	2023	4	28	-155,025.3304	2023-04
2023-10-01	2023	4	29	3,438.8906	2023-04
2023-10-01	2023	4	30	11,804,841.5150	2023-04
2023-10-01	2023	4	31	12,829.6736	2023-04
2023-10-01	2023	4	32	-7,847.4777	2023-04

[99 rows x 5 columns]

Comparing Predicted Sales with Actual Sales

```
df = df.rename(columns={'family_encoded': 'family', 'sales_lag':
'Sales_Lag'})
df_final = df_final.rename(columns={'Prediction': 'Predicted_Sales'})

# Reset index to work with 'quarter_start'
df_reset = df.reset_index()
df_final_reset = df_final.reset_index()

# Group and sum sales by Year-Quarter
df_grouped = df_reset.groupby('Year-Quarter')
['Sales_Lag'].sum().reset_index()
df_final_grouped = df_final_reset.groupby('Year-Quarter')
```

```
['Predicted_Sales'].sum().reset_index()

# Merge on Year-Quarter
merged = pd.merge(df_grouped, df_final_grouped, on='Year-Quarter')

# Sort for plotting
merged = merged.sort_values('Year-Quarter')

# Plotting: scatter + line plot
plt.figure(figsize=(10, 6))

# Actual Sales
plt.plot(merged['Year-Quarter'], merged['Sales_Lag'], linestyle='-',
color='blue', label='Actual Sales (Sales_Lag)')
plt.scatter(merged['Year-Quarter'], merged['Sales_Lag'], color='blue')

# Predicted Sales
plt.plot(merged['Year-Quarter'], merged['Predicted_Sales'],
linestyle='--', color='orange', label='Predicted Sales')
plt.scatter(merged['Year-Quarter'], merged['Predicted_Sales'],
color='orange')

# Decorations
plt.title("Scatter Line Plot: Actual vs Predicted Sales per Quarter")
plt.xlabel("Year-Quarter")
plt.ylabel("Total Sales")
plt.xticks(rotation=45)
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.show()
```



```
# Extract quarter number from column names
final_df['Quarter'] = final_df['Quarter'].str.extract(r'final_Q(\d)').astype(int)

# Add constant year and quarter_start based on your original structure
final_df['Year'] = 2023
quarter_start_map = {2: '2023-04-01', 3: '2023-07-01', 4: '2023-10-01'}
final_df['quarter_start'] = final_df['Quarter'].map(quarter_start_map)

# Reorder columns to match your original df_final structure
final_df = final_df[['quarter_start', 'Year', 'Quarter', 'family', 'Final_Prediction']]

final_df
```

	quarter_start	Year	Quarter	family	Final_Prediction
0	2023-04-01	2023	2	0	-2,479.7427
1	2023-04-01	2023	2	1	1,189.3288
2	2023-04-01	2023	2	2	14,124.9114
3	2023-04-01	2023	2	3	10,209,675.3711
4	2023-04-01	2023	2	4	0.0000
..
94	2023-10-01	2023	4	28	-768,937.0525
95	2023-10-01	2023	4	29	-558.4717
96	2023-10-01	2023	4	30	38,668,590.2551
97	2023-10-01	2023	4	31	60,260.5812
98	2023-10-01	2023	4	32	-23,381.2932

[99 rows x 5 columns]

Replacing Negative Sales from 0

```
# Replace negative Final_Prediction values with 0
final_df.loc[final_df['Final_Prediction'] < 0, 'Final_Prediction'] = 0
final_df
```

	quarter_start	Year	Quarter	family	Final_Prediction
0	2023-04-01	2023	2	0	0.0000
1	2023-04-01	2023	2	1	1,189.3288
2	2023-04-01	2023	2	2	14,124.9114
3	2023-04-01	2023	2	3	10,209,675.3711
4	2023-04-01	2023	2	4	0.0000
..
94	2023-10-01	2023	4	28	0.0000
95	2023-10-01	2023	4	29	0.0000
96	2023-10-01	2023	4	30	38,668,590.2551
97	2023-10-01	2023	4	31	60,260.5812
98	2023-10-01	2023	4	32	0.0000

[99 rows x 5 columns]